

DATA 624 PREDICTIVE ANALYTICS - Project 2

Melissa Bowman, Frederick Jones, Shoshana Farber, Gabriel Campos

Last edited May 05, 2024

Libraries

```
library(Amelia)
library(car)
library(caret)
library(corrplot)
library(Cubist)
library(DataExplorer)
library(dplyr)
library(e1071)
library(earth)
library(forcats)
library(forecast)
library(fpp3)
library(gbm)
library(ggplot2)
library(kableExtra)
library(MASS)
library(mice)
library(mlbench)
library(party)
library(randomForest)
library(RANN)
library(RColorBrewer)
library(readxl)
library(rpart)
library(rpart.plot)
library(summarytools)
library(tidyr)
library(VIM)
library(earth)
library(randomForest)
```

Assignment Description

Project #2 (Team) Assignment

This is role playing. I am your new boss. I am in charge of production at ABC Beverage and you are a team of data scientists reporting to me. My leadership has told me that new regulations are requiring us to

understand our manufacturing process, the predictive factors and be able to report to them our predictive model of pH.

Please use the historical data set I am providing. Build and report the factors in BOTH a technical and non-technical report. I like to use Word and Excel. Please provide your non-technical report in a business friendly readable document and your predictions in an Excel readable format. The technical report should show clearly the models you tested and how you selected your final approach. Please submit both Rpubs links and .rmd files or other readable formats for technical and non-technical reports. Also submit the excel file showing the prediction of your models for pH.

Data Import

We will first load in the data that is required for this analysis.

```
train_df <- readxl::read_xlsx('Data/StudentData.xlsx')
test_df <- readxl::read_xlsx('Data/StudentData.xlsx')
```

StudentData.xlsx is our Training data set.

StudentEvaluation.xlsx is our Test data set.

Exporatory Data Analysis

First, we can preview our dataset.

```
glimpse(train_df)
```

```
## Rows: 2,571
## Columns: 33
## $ `Brand Code`      <chr> "B", "A", "B", "A", "A", "A", "A", "B", "B", "B", ~
## $ `Carb Volume`     <dbl> 5.340000, 5.426667, 5.286667, 5.440000, 5.486667, ~
## $ `Fill Ounces`     <dbl> 23.96667, 24.00667, 24.06000, 24.00667, 24.31333, ~
## $ `PC Volume`       <dbl> 0.2633333, 0.2386667, 0.2633333, 0.2933333, 0.1113~
## $ `Carb Pressure`   <dbl> 68.2, 68.4, 70.8, 63.0, 67.2, 66.6, 64.2, 67.6, 64~
## $ `Carb Temp`       <dbl> 141.2, 139.6, 144.8, 132.6, 136.8, 138.4, 136.8, 1~
## $ PSC               <dbl> 0.104, 0.124, 0.090, NA, 0.026, 0.090, 0.128, 0.15~
## $ `PSC Fill`        <dbl> 0.26, 0.22, 0.34, 0.42, 0.16, 0.24, 0.40, 0.34, 0.~
## $ `PSC CO2`         <dbl> 0.04, 0.04, 0.16, 0.04, 0.12, 0.04, 0.04, 0.04, 0.~
## $ `Mnf Flow`        <dbl> -100, -100, -100, -100, -100, -100, -100, -100, -1~
## $ `Carb Pressure1`  <dbl> 118.8, 121.6, 120.2, 115.2, 118.4, 119.6, 122.2, 1~
## $ `Fill Pressure`   <dbl> 46.0, 46.0, 46.0, 46.4, 45.8, 45.6, 51.8, 46.8, 46~
## $ `Hyd Pressure1`   <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
## $ `Hyd Pressure2`   <dbl> NA, NA, NA, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
## $ `Hyd Pressure3`   <dbl> NA, NA, NA, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
## $ `Hyd Pressure4`   <dbl> 118, 106, 82, 92, 92, 116, 124, 132, 90, 108, 94, ~
## $ `Filler Level`    <dbl> 121.2, 118.6, 120.0, 117.8, 118.6, 120.2, 123.4, 1~
## $ `Filler Speed`    <dbl> 4002, 3986, 4020, 4012, 4010, 4014, NA, 1004, 4014~
## $ Temperature      <dbl> 66.0, 67.6, 67.0, 65.6, 65.6, 66.2, 65.8, 65.2, 65~
## $ `Usage cont`      <dbl> 16.18, 19.90, 17.76, 17.42, 17.68, 23.82, 20.74, 1~
## $ `Carb Flow`       <dbl> 2932, 3144, 2914, 3062, 3054, 2948, 30, 684, 2902,~
## $ Density           <dbl> 0.88, 0.92, 1.58, 1.54, 1.54, 1.52, 0.84, 0.84, 0.~
```

```
## $ MFR <dbl> 725.0, 726.8, 735.0, 730.6, 722.8, 738.8, NA, NA, ~
## $ Balling <dbl> 1.398, 1.498, 3.142, 3.042, 3.042, 2.992, 1.298, 1~
## $ `Pressure Vacuum` <dbl> -4.0, -4.0, -3.8, -4.4, -4.4, -4.4, -4.4, -4~
## $ PH <dbl> 8.36, 8.26, 8.94, 8.24, 8.26, 8.32, 8.40, 8.38, 8.~
## $ `Oxygen Filler` <dbl> 0.022, 0.026, 0.024, 0.030, 0.030, 0.024, 0.066, 0~
## $ `Bowl Setpoint` <dbl> 120, 120, 120, 120, 120, 120, 120, 120, 120, ~
## $ `Pressure Setpoint` <dbl> 46.4, 46.8, 46.6, 46.0, 46.0, 46.0, 46.0, 46~
## $ `Air Pressurer` <dbl> 142.6, 143.0, 142.0, 146.2, 146.2, 146.6, 146.2, 1~
## $ `Alch Rel` <dbl> 6.58, 6.56, 7.66, 7.14, 7.14, 7.16, 6.54, 6.52, 6.~
## $ `Carb Rel` <dbl> 5.32, 5.30, 5.84, 5.42, 5.44, 5.44, 5.38, 5.34, 5.~
## $ `Balling Lvl` <dbl> 1.48, 1.56, 3.28, 3.04, 3.04, 3.02, 1.44, 1.44, 1.~
```

The dataset consists of 2,571 rows and 33 columns. Most of the variables are numeric, except for the first column indicating Brand Code. Our response variable is PH.

We can take also take a look at the summary statistics for each of the numeric variables.

```
summary(train_df)
```

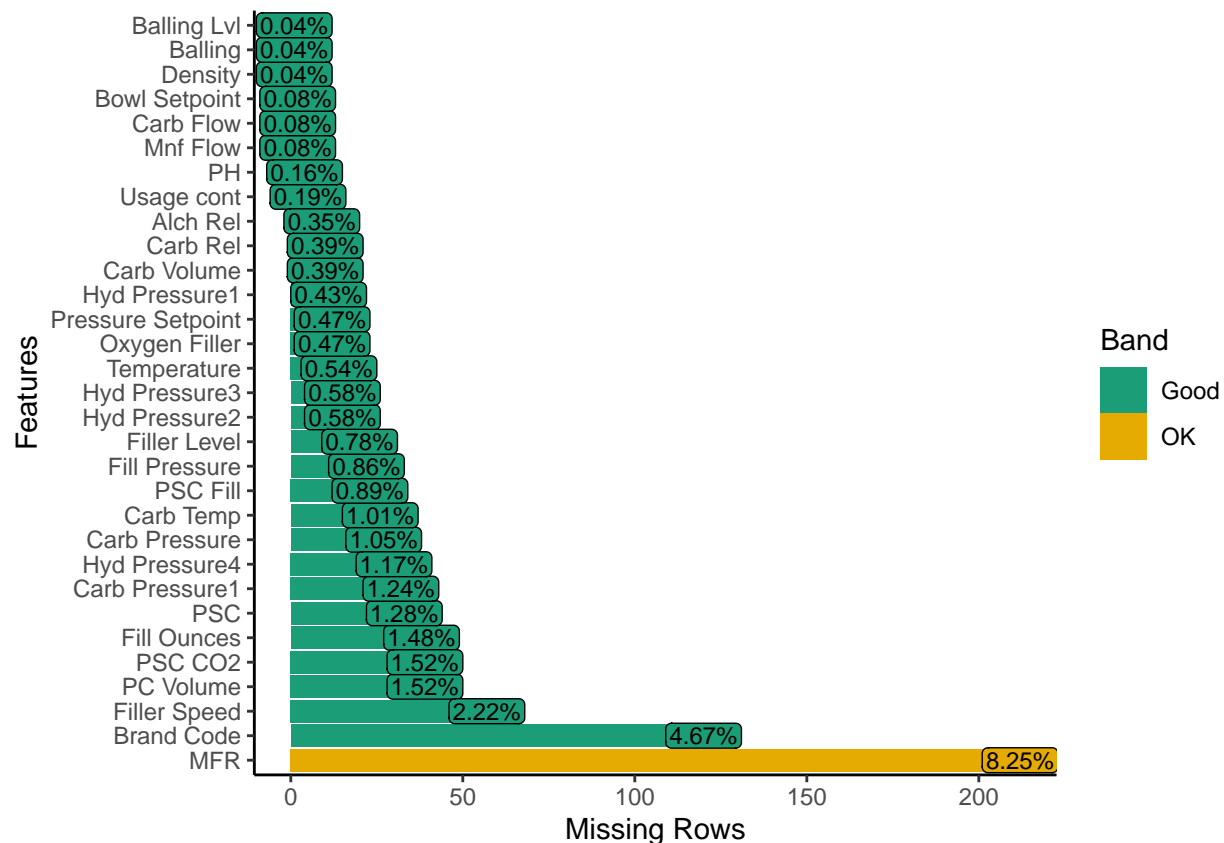
```
## Brand Code Carb Volume Fill Ounces PC Volume
## Length:2571 Min. :5.040 Min. :23.63 Min. :0.07933
## Class :character 1st Qu.:5.293 1st Qu.:23.92 1st Qu.:0.23917
## Mode :character Median :5.347 Median :23.97 Median :0.27133
## Mean :5.370 Mean :23.97 Mean :0.27712
## 3rd Qu.:5.453 3rd Qu.:24.03 3rd Qu.:0.31200
## Max. :5.700 Max. :24.32 Max. :0.47800
## NA's :10 NA's :38 NA's :39
## Carb Pressure Carb Temp PSC PSC Fill
## Min. :57.00 Min. :128.6 Min. :0.00200 Min. :0.0000
## 1st Qu.:65.60 1st Qu.:138.4 1st Qu.:0.04800 1st Qu.:0.1000
## Median :68.20 Median :140.8 Median :0.07600 Median :0.1800
## Mean :68.19 Mean :141.1 Mean :0.08457 Mean :0.1954
## 3rd Qu.:70.60 3rd Qu.:143.8 3rd Qu.:0.11200 3rd Qu.:0.2600
## Max. :79.40 Max. :154.0 Max. :0.27000 Max. :0.6200
## NA's :27 NA's :26 NA's :33 NA's :23
## PSC CO2 Mnf Flow Carb Pressure1 Fill Pressure
## Min. :0.00000 Min. : -100.20 Min. :105.6 Min. :34.60
## 1st Qu.:0.02000 1st Qu.: -100.00 1st Qu.:119.0 1st Qu.:46.00
## Median :0.04000 Median : 65.20 Median :123.2 Median :46.40
## Mean :0.05641 Mean : 24.57 Mean :122.6 Mean :47.92
## 3rd Qu.:0.08000 3rd Qu.:140.80 3rd Qu.:125.4 3rd Qu.:50.00
## Max. :0.24000 Max. : 229.40 Max. :140.2 Max. :60.40
## NA's :39 NA's :2 NA's :32 NA's :22
## Hyd Pressure1 Hyd Pressure2 Hyd Pressure3 Hyd Pressure4
## Min. : -0.80 Min. : 0.00 Min. : -1.20 Min. : 52.00
## 1st Qu.: 0.00 1st Qu.: 0.00 1st Qu.: 0.00 1st Qu.: 86.00
## Median :11.40 Median :28.60 Median :27.60 Median : 96.00
## Mean :12.44 Mean :20.96 Mean :20.46 Mean : 96.29
## 3rd Qu.:20.20 3rd Qu.:34.60 3rd Qu.:33.40 3rd Qu.:102.00
## Max. :58.00 Max. :59.40 Max. :50.00 Max. :142.00
## NA's :11 NA's :15 NA's :15 NA's :30
## Filler Level Filler Speed Temperature Usage cont Carb Flow
## Min. : 55.8 Min. : 998 Min. :63.60 Min. :12.08 Min. : 26
## 1st Qu.: 98.3 1st Qu.:3888 1st Qu.:65.20 1st Qu.:18.36 1st Qu.:1144
```

```
## Median :118.4 Median :3982 Median :65.60 Median :21.79 Median :3028
## Mean :109.3 Mean :3687 Mean :65.97 Mean :20.99 Mean :2468
## 3rd Qu.:120.0 3rd Qu.:3998 3rd Qu.:66.40 3rd Qu.:23.75 3rd Qu.:3186
## Max. :161.2 Max. :4030 Max. :76.20 Max. :25.90 Max. :5104
## NA's :20 NA's :57 NA's :14 NA's :5 NA's :2
## Density MFR Balling Pressure Vacuum
## Min. :0.240 Min. : 31.4 Min. : -0.170 Min. : -6.600
## 1st Qu.:0.900 1st Qu.:706.3 1st Qu.: 1.496 1st Qu.: -5.600
## Median :0.980 Median :724.0 Median : 1.648 Median : -5.400
## Mean :1.174 Mean :704.0 Mean : 2.198 Mean : -5.216
## 3rd Qu.:1.620 3rd Qu.:731.0 3rd Qu.: 3.292 3rd Qu.: -5.000
## Max. :1.920 Max. :868.6 Max. : 4.012 Max. : -3.600
## NA's :1 NA's :212 NA's :1
## PH Oxygen Filler Bowl Setpoint Pressure Setpoint
## Min. :7.880 Min. :0.00240 Min. : 70.0 Min. :44.00
## 1st Qu.:8.440 1st Qu.:0.02200 1st Qu.:100.0 1st Qu.:46.00
## Median :8.540 Median :0.03340 Median :120.0 Median :46.00
## Mean :8.546 Mean :0.04684 Mean :109.3 Mean :47.62
## 3rd Qu.:8.680 3rd Qu.:0.06000 3rd Qu.:120.0 3rd Qu.:50.00
## Max. :9.360 Max. :0.40000 Max. :140.0 Max. :52.00
## NA's :4 NA's :12 NA's :2 NA's :12
## Air Pressurer Alch Rel Carb Rel Balling Lvl
## Min. :140.8 Min. :5.280 Min. :4.960 Min. :0.00
## 1st Qu.:142.2 1st Qu.:6.540 1st Qu.:5.340 1st Qu.:1.38
## Median :142.6 Median :6.560 Median :5.400 Median :1.48
## Mean :142.8 Mean :6.897 Mean :5.437 Mean :2.05
## 3rd Qu.:143.0 3rd Qu.:7.240 3rd Qu.:5.540 3rd Qu.:3.14
## Max. :148.2 Max. :8.620 Max. :6.060 Max. :3.66
## NA's :9 NA's :10 NA's :1
```

NA Proportions

We can plot the missing values for each column to see what proportion of each variable is missing.

```
plot_missing(train_df,
             missing_only = T,
             ggtheme = theme_classic(),
             theme_config = list(legend.position = c("right")),
             geom_label_args = list("size" = 3, "label.padding" = unit(0.1, "lines")))
```




We can see that majority of the variables are missing less than 1% of values. For those that are missing more than 1% of the data, majority still fall below 5%. The variable with the most missing data, and possibly cause for concern, is MFR. However, even this is missing only about 8.25% of the data.

```
data.frame(missing = colSums(is.na(train_df))) |>
  filter(missing == 0) |>
  rownames()
```


```
## [1] "Pressure Vacuum" "Air Pressurer"
```

Pressure Vacuum and Air Pressurer are the only variables not missing any data.

```
VIM::aggr(train_df, numbers=T, sortVars=T, bars = FALSE,
  cex.axis = .6)
```



Combinations



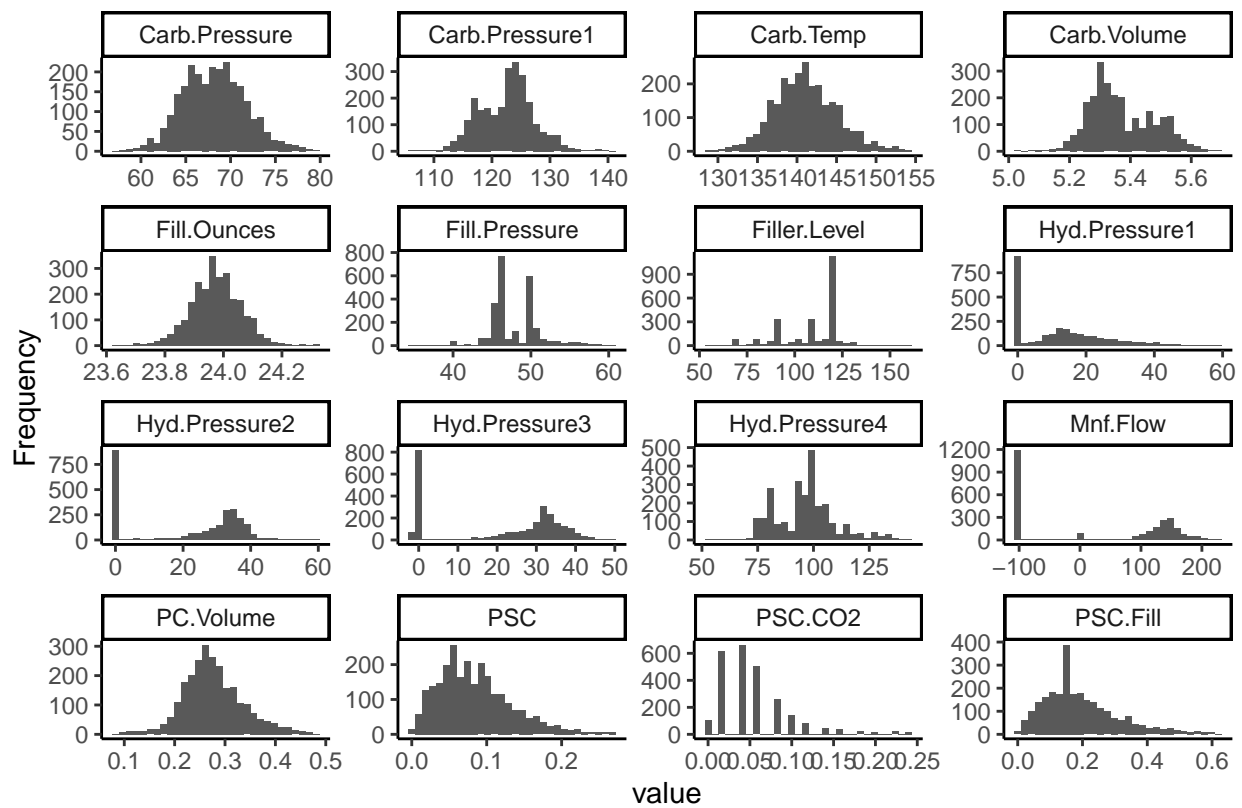
```
##
## Variables sorted by number of missings:
## Variable Count
## MFR 0.0824581875
## Brand Code 0.0466744457
## Filler Speed 0.0221703617
## PC Volume 0.0151691949
## PSC CO2 0.0151691949
## Fill Ounces 0.0147802412
## PSC 0.0128354726
## Carb Pressure1 0.0124465189
## Hyd Pressure4 0.0116686114
## Carb Pressure 0.0105017503
## Carb Temp 0.0101127966
## PSC Fill 0.0089459354
## Fill Pressure 0.0085569817
## Filler Level 0.0077790743
## Hyd Pressure2 0.0058343057
## Hyd Pressure3 0.0058343057
## Temperature 0.0054453520
## Oxygen Filler 0.0046674446
## Pressure Setpoint 0.0046674446
## Hyd Pressure1 0.0042784909
## Carb Volume 0.0038895371
## Carb Rel 0.0038895371
## Alch Rel 0.0035005834
```

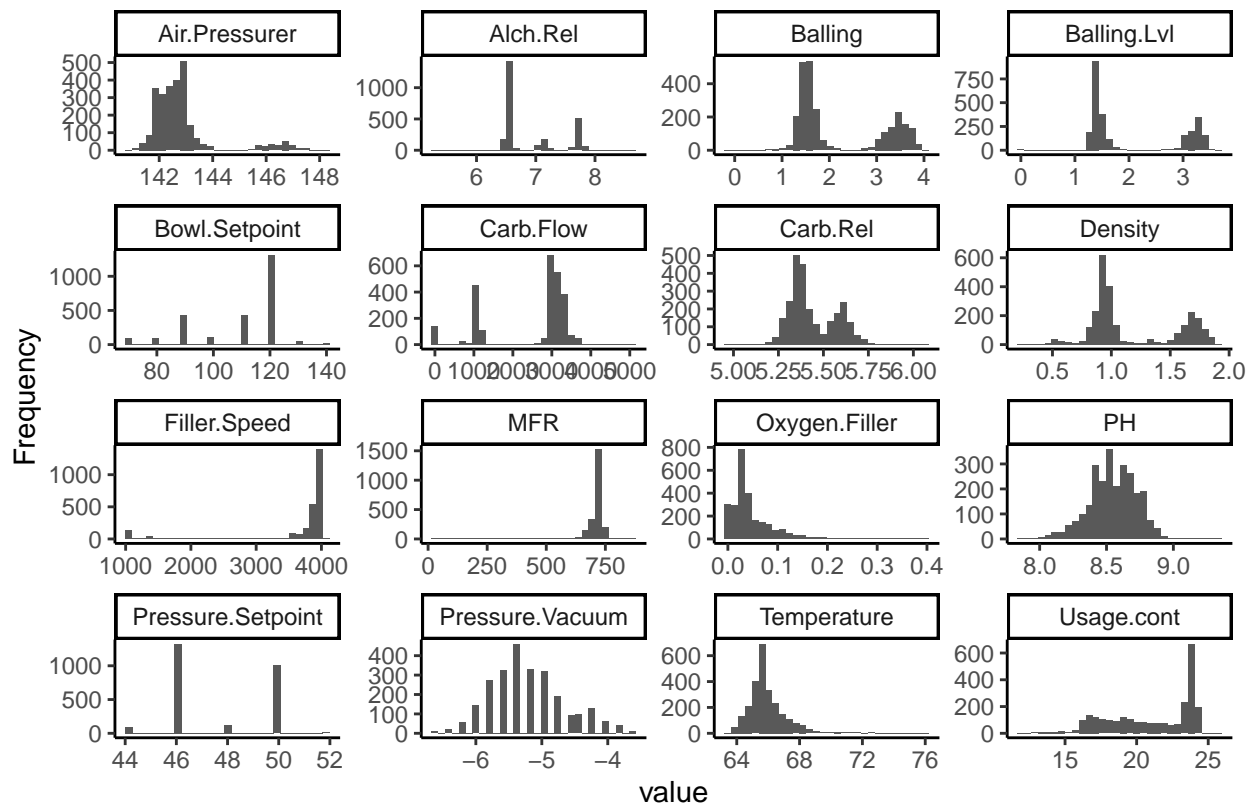
```
##      Usage cont 0.0019447686
##      PH 0.0015558149
##      Mnf Flow 0.0007779074
##      Carb Flow 0.0007779074
##      Bowl Setpoint 0.0007779074
##      Density 0.0003889537
##      Balling 0.0003889537
##      Balling Lvl 0.0003889537
##      Pressure Vacuum 0.0000000000
##      Air Pressurer 0.0000000000
```

Distributions

We will now take a look at the distributions of the numeric variables.

```
DataExplorer::plot_histogram(train_df, nrow = 4L, ncol = 4L, ggtheme = theme_classic())
```





Page 2

Carb Pressure, Carb Temp, Fill Ounces, PC Volume, and PH seem to be relatively normally distributed.

Hyd Pressure 1, PCS, PSC CO2, PSC Fill, Air Pressurer, Oxygen Filler, Pressure Vacuum, and Temperature all seem to have a right skew.

Hyd Pressure2, Hyd Pressure3, and Mnf Flow all seem to have a left skew, although there are also a fair amount of entries with a value at 0. Filler Speed and MFR also seem to have a left skew.

Some variables, such as Balling, Balling Lvl, Carb Rel, and Density seem to be bimodally distributed.

Initial Findings

- Data consists of 2571 observations with 33 columns
- Brand Code:
 - Type character
 - Unordered categorical values
- Predictors:
 - Primarily doubles
 - 4 can be considered integers
 - High range variables:
 - i. Mnf Flow -100.20 to 220.40
 - ii. Hyd Pressure1 -50.00 to 50.00
 - iii. Hyd Pressure2 -50.00 to 61.40
 - iv. Hyd Pressure3 -50.00 to 49.20
 - v. Hyd Pressure4 68.00 to 140.00

- About 8% of the values for MFR is missing.
- Brand Code is missing about 5%
- Filler Speed is missing about 2%
- Remaining Variables have roughly 1% or less missing.
- Pressure.Vacuum, Air.Pressurer have no NAs
- The Distribution of the variables can be grouped as **left skewed**, **right skewed** and for symmetric we can categorized as **relatively normal**
 - Relatively Normal Distributions:
 - * Carb.Pressure
 - * Carb.Temp -Fill.Ounces
 - * PC.Volume
 - * PH
 - Left-skew Distributions:
 - * Carb.Flow
 - * Filler.Speed
 - * Mnf.Flow
 - * MFR
 - * Bowl.Setpoint
 - * Filler.Level
 - * Hyd.Pressure2
 - * Hyd.Pressure3 -Usage.cont
 - * Carb.Pressure1
 - * Filler.Speed
 - Right-skew Distributions:
 - * Pressure.Setpoint
 - * Fill.Pressure
 - * Hyd.Pressure1
 - * Temperature
 - * Carb.Volume
 - * PSC
 - * PSC.CO2
 - * PSC.Fill
 - * Balling
 - * Density
 - * Hyd.Pressure4
 - * Air.Pressurer
 - * Alch.Rel
 - * Carb.Rel
 - * Oxygen.Filler
 - * Balling.Lvl
 - * Pressure.Vacuum


```
unique(train_df$`Brand Code`)
```

```
## [1] "B" "A" "C" "D" NA
```

Brand Code Distribution

Brand Code has 4 categorical values outside of NA (**A,B,C,D**). Let's examine the distribution of these codes.

```
train_df |>
  mutate(`Brand Code` = factor(`Brand Code`, levels = names(sort(table(`Brand Code`), decreasing = TRUE)))
  ggplot(aes(x = `Brand Code`, fill = `Brand Code`)) +
  geom_bar(stat = "count") +
  geom_text(stat = 'count', aes(label = ..count..), vjust = -0.5, color = "black") +
  labs(title = 'Brand Code Distribution', x = 'Brand Code', y = 'Frequency') +
  theme_minimal()
```



Majority of the entries in the dataset belong to **Brand Code B**. A and C have about the same number of entries. There are 120 missing values for **Brand Code**.

Correlation

First, we can plot a correlation matrix of our predictor variables to see which predictors are correlated with each other.

```
train_numeric_df <- train_df |>
  dplyr::select(where(is.numeric)) |>
  na.omit()


# Calculate correlation matrix
train_numeric_cor <- train_numeric_df |>
```

```

dplyr::select(-PH) |>
cor()

# Generate the correlation plot
corrplot(train_numeric_cor,
          method = "color",
          tl.col = "black",
          col = brewer.pal(n = 10,
                           name = "RdYlBu"),
          type = "lower",
          diag=FALSE,
          order = "hclust",
          addCoef.col = "black",
          number.cex = 0.8,
          tl.cex = 0.8,
          cl.cex = 0.8,
          addCoefasPercent = TRUE,
          number.digits = 1)

```



We can see a few instances of multicollinearity in our predictor variables. Carb Rel, Alch Rel, Density, Balling and Balling Level are all significantly positively correlated with each other. Hyd Pressue2 is significantly positively correlated with Hyd Pressure 3. Likewise, Carb Temp with Carb Pressure, MFR with Fill Speed, Bowl Setpoint with Fill Level, and Pressure Setpoint with Fill Pressure.

There are also a number of variables that are highly negatively correlated with each other, such as Pressure Vacuum with Hyd Pressure2 and Hyd Pressure3, Mnf Flow with Filler Level and Bowl Setpoint, and Hyd Pressure4 with Alch Rel.


A number of other variables also display moderate correlations with each other, as can be seen from the medium blue and medium red squares in the correlation plot.

We will need to address these multicollinearity issues in our models.

PH

With PH being our response variable, assessing PH's correlation with other variables is needed.

```
train_numeric_df |>
  dplyr::select(-PH) |> # Exclude 'PH' from predictors if needed
  cor(train_numeric_df$PH) |> # Calculate correlations with 'PH'
  as.data.frame() |>
  rownames_to_column(var = "Predictor") |>
  filter(Predictor != "PH") |> # Ensure 'PH' is not included as its own predictor
  mutate(Predictor = fct_reorder(factor(Predictor), V1)) |> # Reorder factors by correlation for plott
  ggplot(aes(x = Predictor, y = V1, label = round(V1, 2))) +
    geom_col(aes(fill = ifelse(V1 < 0, "negative", "positive"))) +
    geom_text(color = "black", size = 3, vjust = -0.3) +
    coord_flip() +
    labs(title = "Correlations: pH", x = "Predictors", y = "Correlation Coefficient") +
    theme_minimal() +
    theme(legend.position = "none")
```



Individually, there are no variables that are extremely correlated with PH. **Mnf Flow** has the largest correlation of about -0.46. The most significantly positively correlated variables with PH are **Bowl Setpoint** and **Filler Level**. The most significantly negatively correlated variables, other than **Mnf Flow**, are **Usage cont**, **Fill Pressure**, and **Pressure Setpoint**.

Data Cleanup and Pre-Processing

First, to make it easier to reference our variables, let's make each column name snakecase.

```
names(train_df) <- snakecase::to_snake_case(names(train_df))
names(test_df) <- snakecase::to_snake_case(names(test_df))
```

Now, as `ph` is our target variable, we will need to remove any rows that do not have a value for this column.

```
train_df <- train_df |>
  filter(!is.na(ph))
```

We will also transform our `brand_code` variable to categorized factors, replacing any NA value with “Unknown”.

```
train_df <- train_df |>
  dplyr::mutate(brand_code = ifelse(is.na(brand_code), "Unknown", brand_code),
    brand_code = factor(brand_code, levels = c('A', 'B', 'C', 'D', 'Unknown'), ordered = FALSE)


test_df <- test_df |>
  dplyr::mutate(brand_code = ifelse(is.na(brand_code), "Unknown", brand_code),
    brand_code = factor(brand_code, levels = c('A', 'B', 'C', 'D', 'Unknown'), ordered = FALSE)
```

We will identify unhelpful columns in the dataset, such as any variables with zero variance or near zero variance.

```
nzv_df <- nearZeroVar(train_df, saveMetrics= TRUE)
nzv_df <- as.data.frame(nzv_df) %>%
  rownames_to_column(var = "Predictor")

nzv_filtered_df <- nzv_df %>%
  filter(nzv == TRUE)

ggplot(nzv_filtered_df, aes(x = Predictor, y = percentUnique, fill = freqRatio > 0.95)) +
  geom_col(position = "dodge") +
  coord_flip() +
  labs(title = "Near-Zero Variance Predictors",
    x = "Predictors",
    y = "Percentage of Unique Values") +
  theme_minimal()
```



```
print(nzv_filtered_df)
```

```
##      Predictor freqRatio percentUnique zeroVar  nzv
## 1 hyd_pressure_1 31.03704      9.544215  FALSE TRUE
```

hyd_pressure_1 is the only variable with near zero variance. We will not include this variable in our modeling.

Finally, we will pre-process the data for modeling.

The data is in the form of a tibble. For pre-processing using the `preProcess()` function from the `caret` package, we need the data in the form of a dataframe. We will use `as.data.frame()` to do this.

```
train_df <- as.data.frame(train_df)
test_df <- as.data.frame(test_df)
```

We will leverage `caret` package method `preProcess` to transform data using methods: + `knnImpute` - nearest neighbor to impute missing data + `nzv` = remove near-zero values identified above + `corr` = filters out highly correlated values addressing multicollinearity + `center` = subtracts the mean of the predictor's data (again from the data in x) from the predictor values + `scale` = divides by the standard deviation. + `BoxCox` = normalizes data * Use the `predict` function to process the list variables created with `preProcess()` to recreate the dataframe.

```
#remove pH from the train data set in order to only transform the predictors
train_preprocess_df <- train_df |>
```



```

dplyr::select(-c(ph))

preProc_ls <- preProcess(train_preprocess_df, method = c("knnImpute", "nzv", "corr", "center", "scale",
train_preProc_df <- predict(preProc_ls, train_preprocess_df)

# rejoin with pH
train_preProc_df$ph <- train_df$ph

#remove pH from the test data set in order to only transform the predictors
test_preprocess_df <- test_df |>
  dplyr::select(-c(ph))

preProc_ls <- preProcess(test_preprocess_df, method = c("knnImpute", "nzv", "corr", "center", "scale",
test_preProc_df <- predict(preProc_ls, test_preprocess_df)

# rejoin with pH
test_preProc_df$ph <- test_df$ph

```

Let's check that no missing values remain.

```

# verify no NAs remain
colSums(is.na(train_preProc_df))

```

```

##      brand_code      carb_volume      fill_ounces      pc_volume
##           0           0           0           0
## carb_pressure      carb_temp           psc      psc_fill
##           0           0           0           0
##      psc_co_2      mnf_flow carb_pressure_1      fill_pressure
##           0           0           0           0
## hyd_pressure_2 hyd_pressure_4      temperature      usage_cont
##           0           0           0           0
##      carb_flow      mfr      pressure_vacuum      oxygen_filler
##           0           0           0           0
## bowl_setpoint pressure_setpoint      air_pressurer      alch_rel
##           0           0           0           0
##      carb_rel      ph
##           0           0

```

Data Partition

We will split the data into an 80:20 training and validation set.

```

set.seed(1234) # for reproducibility

training_set_df <- createDataPartition(train_preProc_df$ph, p=0.8, list=FALSE)

train <- train_preProc_df[training_set_df,]
eval <- train_preProc_df[-training_set_df,]

```

We will now build several model using the data and we will evaluate each one to determine which is the best model for our data.

Modeling

Linear Model

```
lm <- lm(ph ~ ., data=train)

summary(lm)

##
## Call:
## lm(formula = ph ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.52114 -0.07953  0.01028  0.08828  0.74107
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    8.5013168   0.0111835  760.167 < 2e-16 ***
## brand_codeB     0.0745672   0.0170588   4.371 1.30e-05 ***
## brand_codeC    -0.0736206   0.0187874  -3.919 9.20e-05 ***
## brand_codeD     0.0775571   0.0160734   4.825 1.50e-06 ***
## brand_codeUnknown -0.0068023  0.0217928  -0.312 0.754969
## carb_volume    -0.0109571   0.0084396  -1.298 0.194332
## fill_ounces    -0.0064805   0.0032356  -2.003 0.045325 *
## pc_volume      -0.0068601   0.0036763  -1.866 0.062180 .
## carb_pressure  -0.0009003   0.0122829  -0.073 0.941576
## carb_temp       0.0059642   0.0111378   0.535 0.592368
## psc            -0.0037035   0.0032107  -1.153 0.248842
## psc_fill       -0.0045761   0.0031837  -1.437 0.150770
## psc_co_2       -0.0059852   0.0031667  -1.890 0.058899 .
## mnf_flow       -0.0705019   0.0061363 -11.489 < 2e-16 ***
## carb_pressure_1  0.0355911   0.0037543   9.480 < 2e-16 ***
## fill_pressure   0.0160609   0.0043311   3.708 0.000214 ***
## hyd_pressure_2  0.0193080   0.0046854   4.121 3.93e-05 ***
## hyd_pressure_4 -0.0016373   0.0045338  -0.361 0.718034
## temperature    -0.0137708   0.0034844  -3.952 8.01e-05 ***
## usage_cont     -0.0223799   0.0038886  -5.755 9.97e-09 ***
## carb_flow       0.0127671   0.0040203   3.176 0.001517 **
## mfr            -0.0004474   0.0034511  -0.130 0.896873
## pressure_vacuum -0.0055955   0.0040339  -1.387 0.165563
## oxygen_filler  -0.0082534   0.0042057  -1.962 0.049850 *
## bowl_setpoint   0.0366908   0.0045881   7.997 2.12e-15 ***
## pressure_setpoint -0.0183039  0.0044016  -4.158 3.34e-05 ***
## air_pressurer   0.0008728   0.0032750   0.266 0.789886
## alch_rel        0.0074259   0.0105147   0.706 0.480123
## carb_rel        0.0072810   0.0063010   1.156 0.248008
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1353 on 2026 degrees of freedom
## Multiple R-squared:  0.4023, Adjusted R-squared:  0.394
```

```
## F-statistic: 48.7 on 28 and 2026 DF, p-value: < 2.2e-16
```

The R^2 for this model is 0.394 and there are a number of insignificant variables in the model. Let's use the `step()` function to remove some of the more insignificant variables.


```
lm_update <- step(lm, direction="both", trace=0)
summary(lm_update)
```

```
##
## Call:
## lm(formula = ph ~ brand_code + carb_volume + fill_ounces + pc_volume +
##      carb_temp + psc_fill + psc_co_2 + mnf_flow + carb_pressure_1 +
##      fill_pressure + hyd_pressure_2 + temperature + usage_cont +
##      carb_flow + pressure_vacuum + oxygen_filler + bowl_setpoint +
##      pressure_setpoint + carb_rel, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.52485 -0.07939  0.01110  0.08905  0.75411
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    8.504677   0.009706  876.202 < 2e-16 ***
## brand_codeB     0.066065   0.012472   5.297 1.30e-07 ***
## brand_codeC    -0.081771   0.015138  -5.402 7.37e-08 ***
## brand_codeD     0.086058   0.012101   7.112 1.58e-12 ***
## brand_codeUnknown -0.013530  0.019369  -0.699 0.484940
## carb_volume    -0.012031   0.005914  -2.035 0.042029 *
## fill_ounces    -0.006480   0.003217  -2.015 0.044083 *
## pc_volume      -0.007950   0.003510  -2.265 0.023606 *
## carb_temp       0.005226   0.003088   1.692 0.090794 .
## psc_fill       -0.005203   0.003122  -1.667 0.095730 .
## psc_co_2       -0.006242   0.003149  -1.982 0.047598 *
## mnf_flow       -0.070819   0.006093 -11.623 < 2e-16 ***
## carb_pressure_1  0.035682   0.003717   9.599 < 2e-16 ***
## fill_pressure   0.016100   0.004187   3.846 0.000124 ***
## hyd_pressure_2  0.019547   0.004579   4.269 2.06e-05 ***
## temperature    -0.013855   0.003412  -4.060 5.08e-05 ***
## usage_cont     -0.022723   0.003845  -5.909 4.02e-09 ***
## carb_flow       0.013418   0.003892   3.447 0.000577 ***
## pressure_vacuum -0.005954   0.003943  -1.510 0.131187
## oxygen_filler   -0.008263   0.004164  -1.984 0.047356 *
## bowl_setpoint   0.036745   0.004514   8.140 6.84e-16 ***
## pressure_setpoint -0.018189  0.004356  -4.175 3.10e-05 ***
## carb_rel        0.008844   0.005929   1.492 0.135945
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1351 on 2032 degrees of freedom
## Multiple R-squared:  0.4017, Adjusted R-squared:  0.3952
## F-statistic: 62.01 on 22 and 2032 DF, p-value: < 2.2e-16
```

The R^2 value increased slightly to about 0.395.

Let's check the diagnostic plots for this model.

```
par(mfrow = c(1,2))
plot(lm_update, which = c(1,2))
```



From the residuals vs fitted plot, there does not seem to be any heteroscedasticity, so constant variance is fulfilled. From the QQ-plot, the residuals seem relatively normally distributed although they diverge from the normal line toward the lower end.

Let's evaluate how this model performs on the evaluation data.

```
lm_pred <- predict(lm_update, eval)
(lm_metrics <- postResample(lm_pred, eval$ph))
```

```
##      RMSE Rsquared      MAE
## 0.1305888 0.3934348 0.1043181
```

The evaluation set has an *RMSE* of 0.13 and an R^2 of 0.39.


PLS Model

```
set.seed(2341)

# generate model
```

```
pls_model <- train(ph ~ .,
  data=train,
  method='pls',
  metric='Rsquared',
  tuneLength=10,
  trControl=trainControl(method = "cv", number = 10))

plot(pls_model)
```



```
pls_model

## Partial Least Squares
##
## 2055 samples
## 25 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1849, 1850, 1849, 1849, 1850, 1850, ...
## Resampling results across tuning parameters:
##
##  ncomp  RMSE      Rsquared  MAE
##    1     0.1521396  0.2351118  0.1195055
##    2     0.1426937  0.3274893  0.1115541
##    3     0.1399277  0.3537450  0.1087935
```

```
##      4      0.1382565  0.3696090  0.1080814
##      5      0.1370713  0.3798388  0.1066452
##      6      0.1367173  0.3833915  0.1064445
##      7      0.1364496  0.3857474  0.1060137
##      8      0.1364197  0.3859793  0.1059609
##      9      0.1362762  0.3875012  0.1058786
##     10      0.1363538  0.3868979  0.1059767
##
## Rsquared was used to select the optimal model using the largest value.
## The final value used for the model was ncomp = 9.
```

The optimal number of components for the PLS model was 9, with a corresponding R^2 of about 0.39.

Let's take a look at the most important variables for the PLS model.

```
plot(varImp(pls_model))
```


```
## Warning: package 'pls' was built under R version 4.3.3

##
## Attaching package: 'pls'

## The following object is masked from 'package:corrplot':
##
##      corrplot

## The following object is masked from 'package:caret':
##
##      R2

## The following object is masked from 'package:stats':
##
##      loadings
```



The most important variable is `mnf_flow`.

Let's evaluate how this model performs on the evaluation data.

```
# evaluate model metrics
pls_pred <- predict(pls_model, eval)
(pls_metrics <- postResample(pls_pred, eval$ph))
```

```
##      RMSE Rsquared      MAE
## 0.1297584 0.4008291 0.1034359
```

The evaluation set for the PLS model has a slightly improved R^2 of 0.40.

KNN Model

```
set.seed(613)

knn_model <- train(ph ~ .,
  data = train,
  method = "knn",
  tuneLength = 10)

knn_model
```


```
## k-Nearest Neighbors
```

```
##
## 2055 samples
## 25 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2055, 2055, 2055, 2055, 2055, 2055, ...
## Resampling results across tuning parameters:
##
## k RMSE Rsquared MAE
## 5 0.1390984 0.3987875 0.10224238
## 7 0.1356105 0.4178307 0.10063952
## 9 0.1340402 0.4277480 0.10008633
## 11 0.1334694 0.4317359 0.09995445
## 13 0.1334360 0.4320796 0.10011994
## 15 0.1337366 0.4297865 0.10053509
## 17 0.1341251 0.4271288 0.10093821
## 19 0.1342850 0.4264191 0.10118382
## 21 0.1345970 0.4241648 0.10158822
## 23 0.1350066 0.4208244 0.10207108
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 13.
```

The optimal k was 13, with a corresponding R^2 value of 0.43. This is improved over both the linear and PLS models.

Let's take a look at the most important variables for this model.

```
plot(varImp(knn_model))
```

For the KNN model, `oxygen_filler` is the most important variable and `mnf_flow` is the second most important variable.

Let's evaluate how this model performs on the evaluation data.

```
knn_pred <- predict(knn_model, eval)
(knn_metrics <- postResample(knn_pred, eval$ph))
```

```
##      RMSE  Rsquared      MAE
## 0.12001314 0.49682031 0.09114698
```

The KNN model performs much better than the linear and PLS models, with an R^2 of about 0.5 on the evaluation set.

MARS Model

```
marsGrid <- expand.grid(.degree = 1:2, .nprune = 2:38)

set.seed(613)

mars_model <- train(ph ~ .,
  data = train,
  method = "earth",
  tuneGrid = marsGrid,
```

```

trControl = trainControl(method = "cv"))

mars_model

## Multivariate Adaptive Regression Spline
##
## 2055 samples
## 25 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1848, 1849, 1849, 1850, 1849, 1851, ...
## Resampling results across tuning parameters:
##
## degree nprune RMSE Rsquared MAE
## 1 2 0.1535873 0.2212785 0.11968101
## 1 3 0.1448777 0.3064319 0.11317776
## 1 4 0.1434915 0.3189151 0.11184578
## 1 5 0.1426340 0.3273142 0.11082134
## 1 6 0.1412078 0.3398698 0.10922470
## 1 7 0.1395662 0.3560322 0.10767575
## 1 8 0.1377089 0.3727176 0.10627053
## 1 9 0.1367601 0.3811219 0.10573257
## 1 10 0.1351608 0.3950965 0.10476198
## 1 11 0.1347769 0.3985716 0.10416362
## 1 12 0.1334250 0.4105160 0.10311047
## 1 13 0.1345734 0.4024726 0.10345874
## 1 14 0.1342337 0.4051460 0.10315766
## 1 15 0.1346415 0.4016102 0.10311742
## 1 16 0.1346262 0.4022586 0.10291964
## 1 17 0.1345258 0.4033978 0.10279534
## 1 18 0.1342011 0.4060808 0.10271496
## 1 19 0.1337828 0.4096137 0.10260958
## 1 20 0.1335053 0.4120417 0.10239299
## 1 21 0.1334387 0.4125922 0.10238966
## 1 22 0.1333905 0.4134484 0.10225990
## 1 23 0.1335464 0.4124154 0.10222786
## 1 24 0.1334183 0.4136389 0.10211975
## 1 25 0.1334758 0.4137009 0.10181612
## 1 26 0.1337060 0.4119887 0.10198165
## 1 27 0.1333982 0.4145935 0.10164635
## 1 28 0.1330357 0.4175293 0.10140635
## 1 29 0.1328101 0.4191835 0.10122453
## 1 30 0.1324069 0.4227461 0.10096221
## 1 31 0.1321572 0.4246428 0.10070312
## 1 32 0.1323947 0.4229480 0.10099467
## 1 33 0.1322344 0.4242455 0.10090254
## 1 34 0.1321323 0.4250261 0.10102388
## 1 35 0.1344771 0.4145926 0.10136173
## 1 36 0.1344011 0.4149405 0.10131512
## 1 37 0.1344278 0.4148383 0.10128010
## 1 38 0.1343354 0.4155427 0.10128834
## 2 2 0.1530944 0.2265794 0.11898679

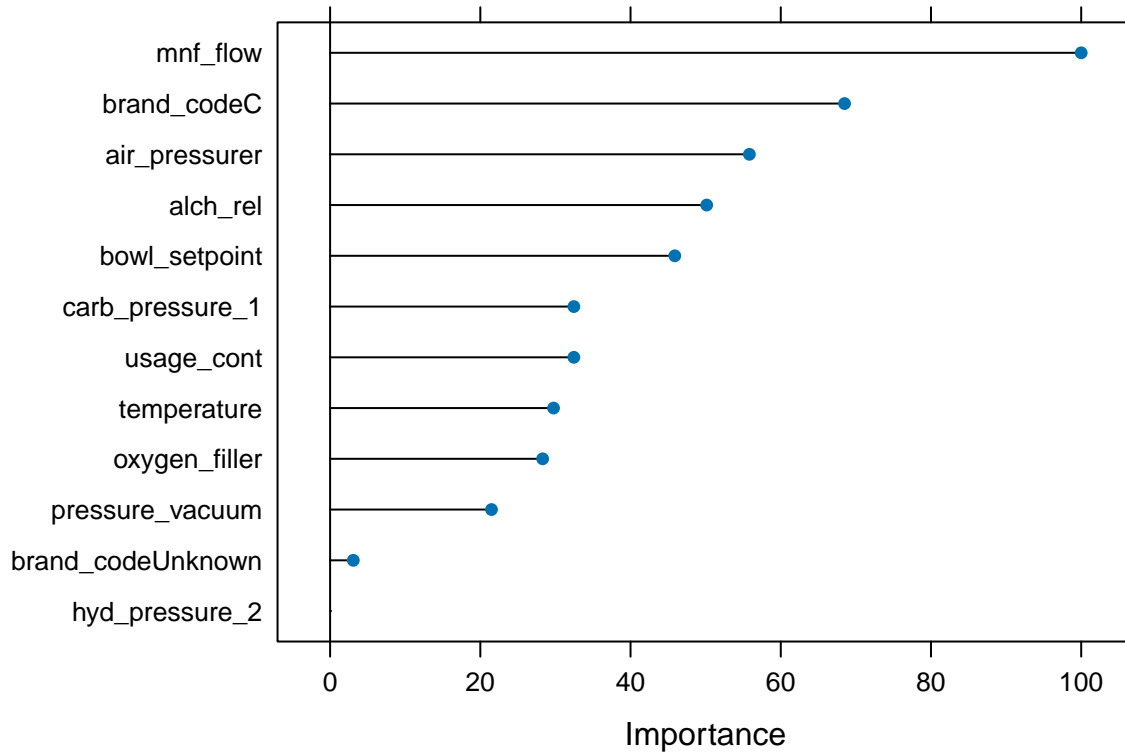
```

```
##      2      3      0.1462250  0.2942891  0.11394808
##      2      4      0.1440784  0.3153016  0.11204350
##      2      5      0.1421114  0.3337770  0.11086995
##      2      6      0.1407611  0.3467498  0.10956907
##      2      7      0.1394805  0.3577464  0.10850651
##      2      8      0.1376499  0.3747515  0.10665540
##      2      9      0.1359799  0.3889583  0.10497439
##      2     10      0.1347076  0.4004098  0.10405077
##      2     11      0.1338458  0.4083320  0.10339104
##      2     12      0.1327088  0.4178045  0.10271164
##      2     13      0.1321305  0.4225395  0.10189380
##      2     14      0.1313585  0.4290817  0.10100493
##      2     15      0.1304205  0.4371764  0.09989178
##      2     16      0.1294247  0.4452763  0.09910754
##      2     17      0.1283584  0.4541348  0.09793470
##      2     18      0.1285692  0.4529985  0.09836378
##      2     19      0.1278478  0.4597029  0.09769820
##      2     20      0.1276282  0.4614473  0.09735486
##      2     21      0.1276065  0.4617771  0.09712116
##      2     22      0.1275896  0.4622986  0.09733336
##      2     23      0.1271724  0.4656356  0.09694150
##      2     24      0.1272758  0.4650817  0.09709139
##      2     25      0.1274095  0.4643671  0.09731865
##      2     26      0.1272583  0.4658037  0.09704997
##      2     27      0.1269681  0.4683432  0.09672071
##      2     28      0.1269133  0.4688745  0.09657392
##      2     29      0.1268942  0.4691397  0.09647809
##      2     30      0.1267669  0.4702629  0.09639925
##      2     31      0.1285202  0.4622434  0.09659245
##      2     32      0.1288632  0.4612861  0.09678814
##      2     33      0.1293031  0.4596399  0.09686214
##      2     34      0.1297509  0.4583308  0.09692435
##      2     35      0.1295061  0.4590350  0.09693595
##      2     36      0.1293917  0.4598075  0.09684188
##      2     37      0.1293917  0.4598075  0.09684188
##      2     38      0.1293917  0.4598075  0.09684188
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were nprune = 30 and degree = 2.
```

The MARS model is optimal at $nprune = 30$ and $degree = 2$. The R^2 at this iteration is 0.47 which is not improved from the KNN model.

Let's take a look at the most important variables for this model.

```
plot(varImp(mars_model))
```



This model has fewer important variables than the PLS and the KNN models. Like with the PLS model, `mnf_flow` is the most important variable.

Let's evaluate how this model performs on the evaluation data.

```
mars_pred <- predict(mars_model, eval)
(mars_metrics <- postResample(mars_pred, eval$ph))
```

```
##      RMSE Rsquared      MAE
## 0.1184498 0.5038378 0.0910705
```

The evaluation set has an R^2 of 0.5, slightly improved over the KNN model.

SVM Model

```
set.seed(613)

svm_model <- train(ph ~ .,
  data = train,
  method = "svmRadial",
  tuneLength = 14,
  trControl = trainControl(method = "cv"))

svm_model
```

```


## Support Vector Machines with Radial Basis Function Kernel
##
## 2055 samples
## 25 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1848, 1849, 1849, 1850, 1849, 1851, ...
## Resampling results across tuning parameters:
##
## C          RMSE          Rsquared    MAE
## 0.25 0.1275992 0.4675130 0.09465438
## 0.50 0.1244643 0.4916487 0.09175646
## 1.00 0.1215873 0.5138680 0.08937335
## 2.00 0.1194783 0.5294227 0.08777912
## 4.00 0.1184050 0.5380039 0.08703835
## 8.00 0.1190210 0.5366143 0.08821722
## 16.00 0.1220096 0.5225359 0.09048471
## 32.00 0.1268173 0.5010522 0.09380972
## 64.00 0.1338263 0.4709833 0.09924286
## 128.00 0.1416055 0.4371453 0.10477848
## 256.00 0.1489243 0.4090902 0.10985776
## 512.00 0.1534970 0.3915444 0.11337869
## 1024.00 0.1535865 0.3916381 0.11343416
## 2048.00 0.1535865 0.3916381 0.11343416
##
## Tuning parameter 'sigma' was held constant at a value of 0.0240063
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.0240063 and C = 4.

```

The optimal model has a sigma of about 0.024 and $C = 4$. The R^2 for this model is about 0.54.

Let's take a look at the most important variables for this model.

```
plot(varImp(svm_model))
```



This model has the same important variables as the MARS model in the same order.

Let's evaluate how this model performs on the evaluation data.

```
svm_pred <- predict(svm_model, eval)
(svm_metrics <- postResample(svm_pred, eval$ph))
```

```
##      RMSE Rsquared      MAE
## 0.1098260 0.5760762 0.0823950
```

This model is much improved from the previous models, with an R^2 of about 0.58 for the evaluation set.

Random Forest Model

```
set.seed(613)

rf_model <- randomForest(ph ~ .,
                          data = train,
                          importance = TRUE,
                          ntree = 1000)

rf_model
```

```
##
```

```
## Call:
## randomForest(formula = ph ~ ., data = train, importance = TRUE, ntree = 1000)
##           Type of random forest: regression
##           Number of trees: 1000
## No. of variables tried at each split: 8
##
##           Mean of squared residuals: 0.01051914
##           % Var explained: 65.14
```

```
rf_model
```

```
##
## Call:
## randomForest(formula = ph ~ ., data = train, importance = TRUE, ntree = 1000)
##           Type of random forest: regression
##           Number of trees: 1000
## No. of variables tried at each split: 8
##
##           Mean of squared residuals: 0.01051914
##           % Var explained: 65.14
```

The model explains 65% of the variability, much improved from our previous models.

Let's take a look at the most important variables for this model.

```
varImp(rf_model) |>
  arrange(desc(Overall)) |>
  knitr::kable()
```

	Overall
brand_code	76.3123199
mnf_flow	57.7016188
pressure_vacuum	52.9055158
usage_cont	50.0934986
oxygen_filler	49.3531816
temperature	42.4104849
alch_rel	41.9715716
air_pressurer	41.5575295
carb_rel	41.1594852
bowl_setpoint	38.8790579
carb_flow	35.6962596
carb_pressure_1	31.7826384
carb_volume	27.5636344
mfr	26.9018896
hyd_pressure_2	21.5836876
pc_volume	21.2972441
fill_pressure	19.8120330
pressure_setpoint	19.6711754
hyd_pressure_4	17.5602769
fill_ounces	7.6233517
carb_pressure	6.5679198
carb_temp	3.9926770

	Overall
psc_co_2	2.0111185
psc_fill	0.2927448
psc	-1.1513903

`brand_code` is the most important variable for this model and `mnf_flow` is the second most important.

Let's evaluate how this model performs on the evaluation data.

```
rf_pred <- predict(rf_model, eval)
(rf_metrics <- postResample(rf_pred, eval$ph))
```

```
##          RMSE  Rsquared          MAE
## 0.09531386 0.68860455 0.07130942
```

This model performs the best from all the previous models. The R^2 for the evaluation set is 0.69.

Let's take a look at all the metrics together.

```
rbind(lm_metrics, pls_metrics, knn_metrics, mars_metrics, svm_metrics, rf_metrics) |>
  knitr::kable()
```

	RMSE	Rsquared	MAE
lm_metrics	0.1305888	0.3934348	0.1043181
pls_metrics	0.1297584	0.4008291	0.1034359
knn_metrics	0.1200131	0.4968203	0.0911470
mars_metrics	0.1184498	0.5038378	0.0910705
svm_metrics	0.1098260	0.5760762	0.0823950
rf_metrics	0.0953139	0.6886045	0.0713094

We can clearly see that the random forest model has the highest prediction accuracy when it comes to the evaluation set, with an R^2 of about 69%.